

---

---

# The Million Song Graph

— Andrea Soto —

---

---

# Project Overview and Goals

Problem	How to store a music recommendation system
Solution	Model the system as a graph to make it easy to transverse relationships
Data	Million Song Dataset and the Last.fm dataset
Project Goals	<ul style="list-style-type: none"><li>- Model data as a graph</li><li>- Implement and automate an ETL process</li><li>- Store data in graph database</li></ul>
Technologies	AWS EC2 and EBS, Jupyter, Spark, Python (pyspark and py2neo libraries), Neo4j

# Motivation

Challenges of implementing a good music recommendation system:

- How to make a good a playlist?
- How to recommend new songs or artists that don't have existing data?
- How to filter songs by things like mood?

→ **Implement a graph model that exploits music interconnectivity**

30 million songs  
2.3 million artists  
1.6 billion playlist



35 million songs



30 million songs



42 million songs



1.5 million songs



# The Million Song Dataset

Collected in 2011 from the Echo Nest API

- 280 GB
- 45,000 unique artists
- 2.2 million artist similarity relations
- HDF5 files (one per song)

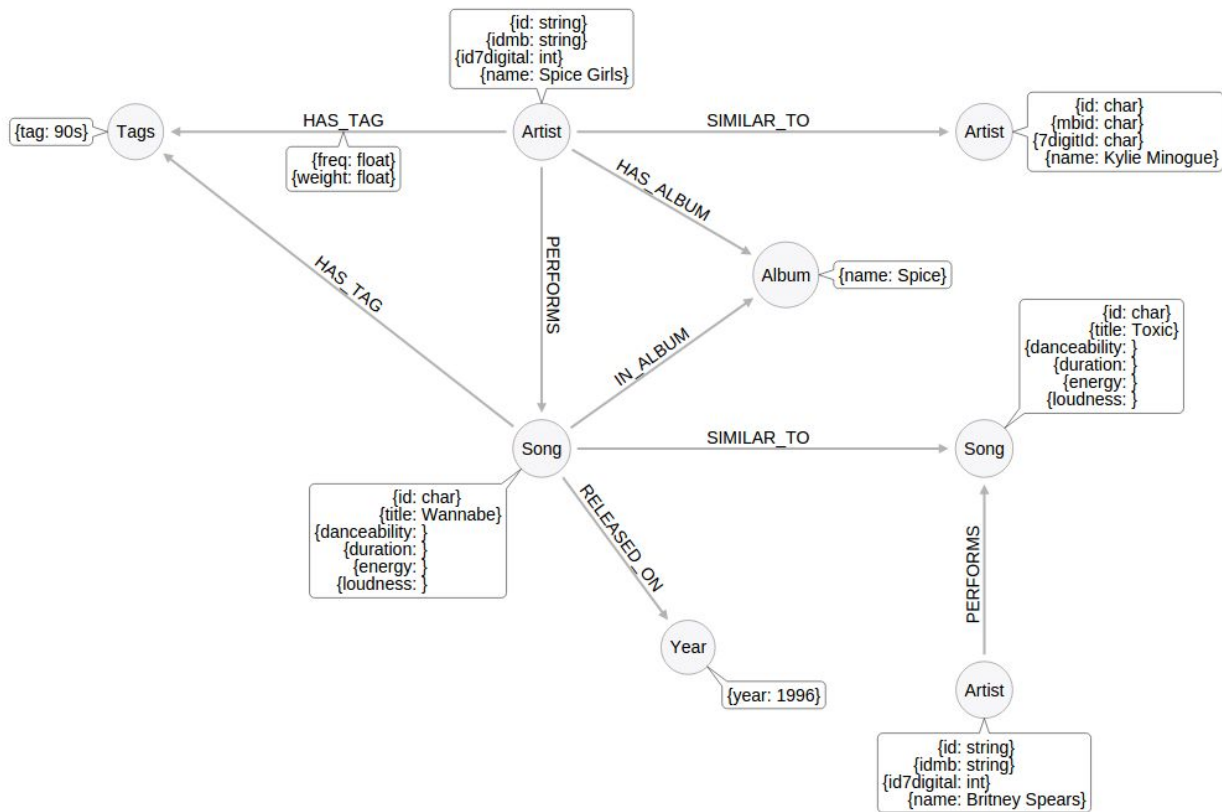
Complemented with last.fm dataset

- 56.5 million song similarity relations
- 8.6 million song-tag pairs
- JSON files



Field Name	Type	Description
artist 7digitalid	int	7digital ID
artist id	string	Echo Nest ID
artist mbid	string	Musicbrainz ID
artist name	string	artist name
artist terms	array string	Echo Nest tags
artist terms freq	array float	Echo Nest tags freqs
artist terms weight	array float	Echo Nest tags weight
danceability	float	algorithmic estimation
duration	float	in seconds
energy	float	from listener point of view
loudness	float	overall loudness in dB
release	string	album name
similar artists	array string	Echo Nest artist IDs
song id	string	Echo Nest song ID
title	string	song title
track id	string	Echo Nest track ID
year	int	MusicBrainz song release

# Graph Model



## Stage 1

10,000 Song Subset  
Direct download of data  
Development and Testing



```
{ 'a_similar': array(['artistId', 'artistId', ..., 'artistId']),  
  'a_terms': array(['term1', 'term2', ..., 'termN']),  
  'a_tfrq': array([ ]),  
  'a_tw': array([ ]),  
  'album': 'album name',  
  'artist_7did': '7digit artist id',  
  'artist_id': 'Echo Nest artist id',  
  'artist_mbid': 'Music Brain artist id',  
  'artist_name': 'Artist name',  
  'dance': 0.0,  
  'dur': 125.7,  
  'energy': 0.0,  
  'loudness': -9.3,  
  'song_id': 'Echo Nest song id',  
  'title': 'Song title',  
  'track_id': 'Echo Nest track id',  
  'year': 1990  
}
```

Spark RDD

## List of files with path

```
[u'./MillionSongSubset/data/B/B/O/TRBBOPX12903D106F7.h5',  
u'./MillionSongSubset/data/B/B/O/TRBBOKQ128F933AE7C.h5',  
u'./MillionSongSubset/data/B/B/O/TRBBOPV12903CFB50F.h5',  
u'./MillionSongSubset/data/B/B/O/TRBBOMJ12903CD18DD.h5',  
u'./MillionSongSubset/data/B/B/O/TRBBOMQ12903CC5186.h5']  
  
[u'./MillionSongSubset/lastfm_subset/B/B/O/TRBBOM128F425DFDC.json',  
u'./MillionSongSubset/lastfm_subset/B/B/O/TRBBOPX12903D106F7.json',  
u'./MillionSongSubset/lastfm_subset/B/B/O/TRBBOMQ12903CC5186.json',  
u'./MillionSongSubset/lastfm_subset/B/B/O/TRBBOME12903CC3862.json',  
u'./MillionSongSubset/lastfm_subset/B/B/O/TRBBOFH128F14A2A46.json']
```

## Stage 2

1,000,000 Dataset  
Attach AWS volume with data

## Parse & Save as CSV

No.	Node Label	File Name	CSV Format
1	Artists	nodes_artists.csv	'artist_id', 'artist_mbid', 'artist_7did', 'artist_name'
2	Songs	nodes_songs.csv	'song_id', 'track_id', 'title', 'dance', 'dur', 'energy', 'loudness'
3	Albums	nodes_albums.csv	'album_name'
4	Year	nodes_years.csv	'year'
5	Tags	nodes_tags.csv	'tag'

No.	Relationship Label	File Name	CSV Format
1	(Artist) - [Similar_To] -> (Artist)	rel_similar_artists.csv	'from_artist_id', 'to_artist_id'
2	(Artist) - [Performs] -> (Song)	rel_performs.csv	'artist_id', 'song_id'
3	(Artist) - [Has_Album] -> (Album)	rel_artist_has_album.csv	'artist_id', 'album_name'
4	(Artist) - [Has_Tag] -> (Tag)	rel_artist_has_tag.csv	'artist_id', 'tag_name', 'frq', 'weight'
5	(Song) - [In_Album] -> (Album)	rel_song_in_album.csv	'song_id', 'album_name'
6	(Song) - [Similar_To] -> (Song)	rel_similar_songs.csv	'from_song_id', 'to_song_id', 'weight'
7	(Song) - [Has_Tag] -> (Tag)	rel_song_has_tag.csv	'from_song_id', 'to_song_id', 'weight'
8	(Song) - [Released_On] -> (Year)	rel_song_year.csv	'song_id', 'year'

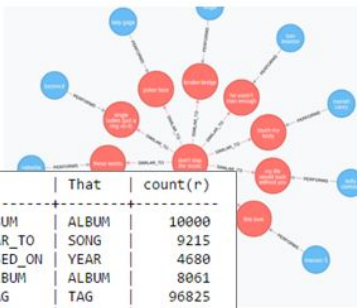
## Batch import



py2neo 2.0

	LABELS(n)	count(n)
1	[u'ALBUM']	7823
2	[u'TAG']	35112
3	[u'SONG']	10000
4	[u'YEAR']	69
5	[u'ARTIST']	3888

	This	To	That	count(r)
1	SONG	IN_ALBUM	ALBUM	10000
2	SONG	SIMILAR_TO	SONG	9215
3	SONG	RELEASED_ON	YEAR	4680
4	ARTIST	HAS_ALBUM	ALBUM	8061
5	ARTIST	HAS_TAG	TAG	96825
6	SONG	HAS_TAG	TAG	99296
7	ARTIST	PERFORMS	SONG	10000
8	ARTIST	SIMILAR_TO	ARTIST	42970



# Live Demo of Music Graph in Neo4j

The screenshot displays the Neo4j 2.3.0 web interface in a browser window. The browser's address bar shows the URL `ec2-54-91-193-97.compute-1.amazonaws.com:7474/browser/`. The interface is divided into a dark sidebar on the left and a main content area on the right.

**Neo4j 2.3.0 Sidebar:**

- Node labels:** Includes buttons for `ALBUM`, `ARTIST`, `SONG`, `TAG`, and `YEAR`.
- Relationship types:** Includes buttons for `HAS_ALBUM`, `HAS_TAG`, `IN_ALBUM`, `PERFORMS`, `RELEASED_ON`, and `SIMILAR_TO`.
- Property keys:** Includes buttons for `danceability`, `duration`, `energy`, `lrq`, `id`, `id7d`, `idmb`, `loudness`, `name`, `trackid`, `tag`, `title`, `trackid`, `weight`, and `year`.
- Database:** Shows the location as `/data/neo4j/data/graph.db` and the size as `142.19 MiB`.

**Main Content Area:**

- At the top, there is a text input field containing `:play start`.
- Below the input field is the Neo4j logo and the text `2.3.0 - COMMUNITY`.
- There are three main sections with icons and buttons:
  - Learn about Neo4j:** A graph epiphany awaits you. Questions include "What is a graph database?", "How can I query a graph?", and "What do people do with Neo4j?". A button labeled `Start Learning` is at the bottom.
  - Jump into code:** Use Cypher, the graph query language. Options include "Code walk-throughs", "RDBMS to Graph", and "Query templates". A button labeled `Write Code` is at the bottom.
  - Monitor the system:** Key system health and status metrics. Options include "Disk utilization", "Cache activity", and "Cluster health and status". A button labeled `Monitor` is at the bottom.
- At the bottom of the main content area, there is a link: [Latest Blog: How Walmart Uses Neo4j for Retail Competitive Advantage - Neo4j Graph Database](#).
- The footer of the interface states: `Copyright © Neo Technology 2002–2015`.

# Pros and Cons of a Graph Database

- Good fit for highly connected data where there tends to be more relationships than nodes
- Powerful for traversal queries (friends-of-friends)
- Cypher is an intuitive and easy to use declarative query language
- Neo4j High Availability enables horizontal scaling of reads, and also supports ACID transactions
- Scales vertically because graph structure is hard to partition (bigger data, bigger server)
- Finding all nodes of the same type is more expensive than in a relational database



# Future Work

- Build a sophisticated algorithm in the back-end that determines the similarity edges in the graph
- Run personalized Page Rank to make recommendations to users
- Add new songs, node types, or more properties by connecting to the APIs of Echo Nest, Musicbrainz, 7digital, or other music APIs
- Add nuances like mood tags, artist collaboration, song covers, live versions, etc.

# Takeaways

- Prototype for migrating music data into Neo4j
- ETL sounds really simple, but can be really complex. Both the data and the technologies used will drive the ETL process
- When handling a larger dataset, thinking about DAGs and the overall process makes a huge difference
- It's also important to understand the limitations of each technology, and know how to tune configurations to improve performance

# Questions?

---